

Pierre Beckmann

pierrebeckmann@gmail.com · pierrebeckmann.com · Google Scholar · ORCID

Education

- May 2025 – Present **EPFL and Idiap Research Institute** – Lausanne and Martigny
PhD student, Neuro-Symbolic AI group of Prof. André Freitas
M-RATIONAL project.
- 2020 – 2023 **University of Lausanne**
M.A. in Philosophy, with 40 catch-up credits
Philosophy of science, Kantian philosophy and phenomenology. Thesis: *Le transcendantal, le physiologique et le computationnel: le predictive processing à partir de Kant et de Helmholtz.*
- 2022 **Humboldt University Berlin**
Exchange semester in philosophy
- 2019 – 2020 **University of Lausanne**
Propaedeutic year in philosophy, sociology and French
- 2016 – 2019 **ETH Zurich**
M.Sc. in Computational Science and Engineering
Machine learning, NLP, computer vision and HPC. Thesis: *Story understanding and story generation with deep learning.*
- 2013 – 2016 **EPFL** – Lausanne
B.Sc. in Life Sciences and Technologies

Research experience

- Jan. 2026 – Present **MATS Scholar** – Machine Learning Alignment & Theory Scholars, Berkeley
Supervised by Patrick Butlin. Extension phase May – Nov. 2026.
- 2023 – 2024 **Doctoral researcher in philosophy of AI** – University of Bern
Research in the philosophy of deep learning, supervised by Prof. Claus Beisbart.
- 2019 – 2021 **Research engineer (50%)** – EPFL LSIR
Implemented and deployed multimodal-LLM-based video retrieval systems.
- 2019 – 2023 **Scientific advisor (10%)** – Logitech, Lausanne
Consulting for internship and master thesis projects on deep learning for audio.
- 2019 **Research engineering intern** – Logitech, Lausanne
Product-feasible deep learning solutions for audio tasks.
- 2018 – 2019 **Master thesis and semester project** – Disney Research Zurich
Story understanding and generation with deep learning.

Teaching experience

- 2024 **Assistant lecturer, Philosophy of AI** – University of Bern
Co-taught two CAS courses in Advanced Machine Learning, delivering one third of the lectures: *Philosophy and Ethics of Extended Cognition and Artificial Intelligence*; *NLP: Philosophical and Ethical Aspects*.
- 2023 **Tutor, Philosophy of Mind** – University of Lausanne
Introductory course in contemporary philosophy of mind for first-year students.
- 2021 – 2022 **Tutor, General and Systematic Philosophy** – University of Lausanne
Introductory course in phenomenology for first-year students; designed and ran three methodology workshops.

Publications in philosophy

A Polyphonic Conception of AI Understanding

Matthieu Queloz, **Pierre Beckmann**.
arXiv, 2026. Paper.

Where is the Mind? Persona Vectors and LLM Individuation

Pierre Beckmann, Patrick Butlin.
Ergo, forthcoming. Paper.

Mechanistic Indicators of Understanding in Large Language Models

Pierre Beckmann, Matthieu Quelo.

Philosophical Studies, 2026. Paper.

Why We Care About Understanding: Competence through Predictive Compression

Matthieu Quelo., **Pierre Beckmann**.

PhilPapers preprint, 2025.

New Horizons in Machine Understanding: Explanatory and Objectual Understanding in Deep Learning Video Generation Models

Pierre Beckmann.

Synthese, 2025. Paper.

An alternative to cognitivism: computational phenomenology and deep learning

Pierre Beckmann, Guillaume Köstner, Ines Hipolito.

Minds and Machines, 2023. Paper.

Publications in LLM research

World Modeling in Transformers

Pierre Beckmann, Matthieu Quelo., André Freitas.

arXiv, 2026. Paper, interactive explorer.

SciR: A Controllable Benchmark for Scientific Reasoning in LLMs

Pierre Beckmann, Marco Valentino, André Freitas.

NeurIPS, 2026. Paper, code, data.

Probing Persona-Dependent Preferences in Language Models

Oscar Gilg, **Pierre Beckmann**, Daniel Paleka, Patrick Butlin.

NeurIPS, 2026. Paper.

Adaptive LLM-Symbolic Reasoning via Dynamic Logical Solver Composition

Lei Xu*, **Pierre Beckmann***, Marco Valentino, André Freitas.

EACL, 2026. Paper.

Publications in deep learning for speech

No longer active in this field.

BYOL-S: Learning self-supervised speech representations by bootstrapping

Gasser Elbanna, Neil Scheidwasser-Clow, Mikolaj Kegler, **Pierre Beckmann**, Karl El Hajal, Milos Cernak.

PMLR, 2023; 3rd place in the HEAR benchmark. Paper, code.

Hybrid Handcrafted and Learnable Audio Representation for Analysis of Speech Under Cognitive and Physical Load

Gasser Elbanna, Alice Biryukov, Neil Scheidwasser-Clow, Lara Orlandic, Pablo Mainar, Mikolaj Kegler, **Pierre Beckmann**, Milos Cernak.

Interspeech, 2022. Paper, code.

SERAB: A multi-lingual benchmark for speech emotion recognition

Neil Scheidwasser-Clow, Mikolaj Kegler, **Pierre Beckmann**, Milos Cernak.

ICASSP, 2021. Paper, code.

Word-Level Embeddings for Cross-Task Transfer Learning in Speech Processing

Pierre Beckmann*, Mikolaj Kegler*, Milos Cernak.

EUSIPCO, 2021. Paper, code.

Deep speech inpainting of time-frequency masks

Pierre Beckmann*, Mikolaj Kegler*, Milos Cernak.

Interspeech, 2020. Paper, demo.

Course book

Neural Networks for Cognitive Sciences

2025. Three chapters, with Prof. Jeff Yoshimi, University of California Merced: history of neural networks; convolutional neural networks; transformers and LLMs. Book.

Talks

Nov. 2026. Talk at The Age of the Bookworm (upcoming)

University of Geneva

Oct. 2026. A map of Manhattan inside of TaxiGPT
Idiap Research Institute, Martigny

July 2026. Mechanistic indicators of understanding in LLMs
Cultural Analytics and Digital Approaches in Philosophy, online seminar

June 2026. Building a reflective equilibrium machine?
Workshop, University of Bern

June 2026. Faut-il se soucier du bien-être des IAs?
Kiosque de Philosophie, Groupe Vaudois de Philosophie, Lausanne

Jan. 2026. Do LLMs understand? A mechanistic perspective
Invited talk, CIS Brownbag Speaker Series, University of California, Merced

Nov. 2025. Why we care about understanding: competence via predictive compression
Philosophy of science colloquium, University of Bern

Oct. 2025. Information recall in deep learning: beyond the feature combination paradigm
PhAI 2025, Vrije Universiteit Amsterdam

Feb. 2025. Computational phenomenology and deep learning
Calculating Experience? seminar, St Catherine's College, Oxford

Nov. 2024. Les neurosciences computationnelles à l'épreuve de la méthode transcendantale
Kant tercentenary series, University of Lausanne

Oct. 2024. Could deep learning video generation models understand the physical world?
Idiap Research Institute

Nov. 2022. Computational phenomenology and deep learning
Series on AI and consciousness, University of Lausanne

Reviewing

Journals

Ergo, Philosophical Studies, Synthese, Res Philosophica, Journal of Consciousness Studies, Knowledge in Society.

Workshops and other

Mechanistic Interpretability Workshop (ICML 2026); MathNLP Workshop (EMNLP 2025, 2026); Digital Minds Research Sprint (2026); Ubiquity Press, *Understanding Conversational AI* (2025).

Grants

2026	AI Safety Support: stipend of USD 36'000 for the MATS extension
2026	AI Safety Support: stipend of USD 14'400 for the MATS program
2023	University of Bern: seed funding of CHF 10'000 for a research period

Service

2024 – Present	Committee member, Groupe Vaudois de Philosophie
----------------	---